

perf工具及PAPI的使用

邹永浩

2019211168

perf 分析热点

因为我是自己编译的内核，所以需要进入内核目录下 `tools/perf` 进行编译安装，直接 `make` 再 `make install` 就可以了。

首先对程序运行概况进行分析，使用 `perf stat` 查看：

```
bugscan@bugscan-laptop:~/Desktop$ perf stat ./a.out
0 25 20 15 0.24 7.03 28946421

Performance counter stats for './a.out':

      7,270.75 msec task-clock:u          #    1.000 CPUs utilized
           0      context-switches:u      #    0.000 K/sec
           0      cpu-migrations:u        #    0.000 K/sec
       4,913      page-faults:u           #   675.791 M/sec
  21,031,111,961 cycles:u                 # 2892862.718 GHz
 34,642,519,923 instructions:u           #    1.65  insn per cycle
 1,567,000,841   branches:u              # 215543444.429 M/sec
 11,937,794      branch-misses:u         #    0.76% of all branches

      7.271439869 seconds time elapsed

      7.259089000 seconds user
      0.011998000 seconds sys
```

可以看到程序是纯计算类的，因为 `task-clock` 为1.

进一步分析程序热点，使用 `perf record` 记录，再使用 `perf report` 查看结果：

```
Samples: 36K of event 'cycles:uppp', Event count (approx.): 26444281452
Overhead Command Shared Object Symbol
95.70% a.out a.out [.] do_main_sieve
1.68% a.out a.out [.] count_zero_bits
1.39% a.out a.out [.] init_main_sieve
1.19% a.out a.out [.] update_small_sieve
0.03% a.out [unknown] [k] 0xffffffff99600a27
0.00% a.out libc-2.27.so [.] _IO_file_xsputn@GLIBC_2.2.5
0.00% a.out ld-2.27.so [.] strcmp
0.00% a.out ld-2.27.so [.] _start
```

可以看到，程序热点函数为 `do_main_sieve`，如果需要进行优化，应重点关注该函数。

PAPI 采集指标

安装 `PAPI` 比较简单，只要把源码下载编译安装即可，编译完成后，我们可以通过 `utils/papi_avail` 查看支持的事件，如下图所示：

```
=====
PAPI Preset Events
=====
Name      Code Avail Deriv Description (Note)
PAPI_L1_DCM 0x80000000 Yes No Level 1 data cache misses
PAPI_L1_ICM 0x80000001 Yes No Level 1 instruction cache misses
PAPI_L2_DCM 0x80000002 Yes Yes Level 2 data cache misses
PAPI_L2_ICM 0x80000003 Yes No Level 2 instruction cache misses
PAPI_L3_DCM 0x80000004 No No Level 3 data cache misses
PAPI_L3_ICM 0x80000005 No No Level 3 instruction cache misses
PAPI_L1_TCM 0x80000006 Yes Yes Level 1 cache misses
PAPI_L2_TCM 0x80000007 Yes No Level 2 cache misses
PAPI_L3_TCM 0x80000008 Yes No Level 3 cache misses
```

可以看到 PAPI 支持查看各类缓存的数据或者指令miss，在我的机器上, 有很多指标无法统计, 其中影响实验的是 PAPI_L1_TCA 以及 PAPI_TLB_TL , 因此我在测试 L1 cache 时仅统计了miss的次数, 对于 TLB 的 miss rate 则未统计.

修改被测程序，添加如下代码：

```
// ...
#include <papi.h>
// ...
int main(int argc, char **argv)
{
    /* Initialize the PAPI library */
    if (PAPI_library_init(PAPI_VER_CURRENT) != PAPI_VER_CURRENT)
        exit(1);
    /* Create an EventSet */
    int EventSet = PAPI_NULL;
    int retval = PAPI_create_eventset(&EventSet);
    assert(retval == PAPI_OK);

    /* Total cycles */
    retval = PAPI_add_event(EventSet, PAPI_TOT_CYC);
    assert(retval == PAPI_OK);
    /* Instructions completed */
    retval = PAPI_add_event(EventSet, PAPI_TOT_INS);
    assert(retval == PAPI_OK);

    /* L1 cache miss */
    retval = PAPI_add_event(EventSet, PAPI_L1_TCM);
    assert(retval == PAPI_OK);
    /* L1 cache access */
    // retval = PAPI_add_event(EventSet, PAPI_L1_TCA);
    // assert(retval == PAPI_OK);
    /* L2 cache miss */
    retval = PAPI_add_event(EventSet, PAPI_L2_TCM);
    assert(retval == PAPI_OK);
    /* L2 cache access */
    retval = PAPI_add_event(EventSet, PAPI_L2_TCA);
    assert(retval == PAPI_OK);
    /* L3 cache miss */
    retval = PAPI_add_event(EventSet, PAPI_L3_TCM);
    assert(retval == PAPI_OK);
    /* L3 cache access */
    retval = PAPI_add_event(EventSet, PAPI_L3_TCA);
    assert(retval == PAPI_OK);

    /* TLB miss */
    // retval = PAPI_add_event(EventSet, PAPI_TLB_TL);
    // assert(retval == PAPI_OK);

    /* branch mispredict */
    retval = PAPI_add_event(EventSet, PAPI_BR_MSP);
    assert(retval == PAPI_OK);
    /* branch */
    retval = PAPI_add_event(EventSet, PAPI_BR_CN);
    assert(retval == PAPI_OK);

    /* Start counting events */
}
```

```

if (PAPI_start(EventSet) != PAPI_OK)
    retval = PAPI_start(EventSet);
assert(retval == PAPI_OK);

long long values1[9];
long long values2[9];
PAPI_read(EventSet, values1);
assert(retval == PAPI_OK);

// logic ignored

/* Stop counting events */
retval = PAPI_stop(EventSet, values2);
assert(retval == PAPI_OK);

int index = 0;
printf("CPI: %f\n", (values2[index] - values1[index]) / 1.0 / (values2[index + 1] - values1[index + 1]));
index += 2;
printf("L1 cache miss: %lld\n", (values2[index] - values1[index]));
index++;
printf("L2 cache miss rate: %f\n", (values2[index] - values1[index]) / 1.0 / (values2[index + 1] - values1[index + 1]));
index += 2;
printf("L3 cache miss rate: %f\n", (values2[index] - values1[index]) / 1.0 / (values2[index + 1] - values1[index + 1]));

// printf("TLB cache miss: %lld\n", (values2[8] - values1[8]));

index += 2;
printf("Branch miss prediction rate: %f\n", (values2[index] - values1[index]) / 1.0 / (values2[index + 1] - values1[index + 1]));

/* Clean up EventSet */
if (PAPI_cleanup_eventset(EventSet) != PAPI_OK)
    exit(-1);

/* Destroy the EventSet */
if (PAPI_destroy_eventset(&EventSet) != PAPI_OK)
    exit(-1);

/* Shutdown PAPI */
PAPI_shutdown();

return 0;
}

```

编译命令如下:

```
gcc prime_sieve-64.c /usr/local/lib/libpapi.a
```

运行结果为:

```
zyh@pc:~/Desktop/MeasuringComputerPerformance/hw11$ ./a.out
0 25 20 15 0.19 4.98 28946421
CPI: 0.565563
L1 cache miss: 614125883
L2 cache miss rate: 1.086262
L3 cache miss rate: 0.091261
Branch miss prediction rate: 0.007873
```

参考文献

<https://www.ibm.com/developerworks/cn/linux/l-cn-perf1/index.html>

<https://www.cnblogs.com/dmyu/p/4648413.html>